

June 10 API Results: GPT-4o vs. Claude Sonnet

1 Data Source

This report uses the frozen pre-June-11 Spanish-time snapshot in `results/api_models_pre_2026_06_11_es`. The snapshot was extracted from commit `96e05c3` (*Llava finetuning fix*), dated 2026-06-10 22:37:44 +0200.

The available Claude file in that snapshot is named `claude-sonnet-4-5`. The comparison below uses that stored result set as the Sonnet side of the requested GPT-4o versus Sonnet analysis.

2 Overall Matched Comparison

The overall comparison uses the 30 benchmarks present for both `gpt-4o` and `claude-sonnet-4-5`. GPT-4o has the higher aggregate score, higher normalized score, and wins more benchmarks head-to-head.

Table 1: Overall matched summary across 30 shared benchmarks.

Model	Benchmarks	Mean score	Median score	Mean normalized
<code>gpt-4o</code>	30	0.428	0.327	0.733
<code>claude-sonnet-4-5</code>	30	0.368	0.225	0.367

Table 2: Pairwise result on the same 30 benchmarks.

Model A	Model B	Wins A	Wins B	Ties	Mean delta A-B
<code>gpt-4o</code>	<code>claude-sonnet-4-5</code>	19	8	3	0.061




Figure 1: Mean benchmark score across all shared June 10 benchmarks.




Figure 2: Mean normalized score across all shared June 10 benchmarks.




Figure 3: Number of first-place finishes in the two-model comparison.




Figure 4: Pairwise win-rate heatmap for GPT-4o and Sonnet.




Figure 5: Distribution of benchmark scores by model.




Figure 6: Bootstrap confidence intervals for mean benchmark score.




Figure 7: Raw benchmark score heatmap.




Figure 8: Normalized benchmark score heatmap.






Figure 9: Benchmark rank heatmap.






Figure 12: Family radar plot.




Figure 13: Family small multiples.

3 Focus on Type A and Type B Benchmarks

Type A contains visual question answering and visual reasoning benchmarks. Type B contains detection and localization benchmarks. The June 10 snapshot has six shared Type A benchmarks and six shared Type B benchmarks for this pair.


Table 3: Type A/B matched summary across 12 shared benchmarks.

Model	Benchmarks	Mean score	Median score	Mean normalized
gpt-4o	12	0.450	0.458	0.833
claude-sonnet-4-5	12	0.322	0.293	0.250

Table 4: Type A/B family means.

Family	Model	Benchmarks	Mean score	Mean rank
Type A / QA	gpt-4o	6	0.800	1.00
Type A / QA	claude-sonnet-4-5	6	0.508	2.00
Type B / Detection	gpt-4o	6	0.100	1.33
Type B / Detection	claude-sonnet-4-5	6	0.135	1.50

GPT-4o is clearly ahead on Type A QA benchmarks. Type B is more mixed: Claude has a higher raw detection mean, while GPT-4o has the better average rank and normalized score because the wins are distributed differently across the six detection benchmarks.

**Figure 14:** Mean score on Type A and Type B benchmarks.




Figure 15: Mean normalized score on Type A and Type B benchmarks.




Figure 16: First-place finishes on Type A and Type B benchmarks.




Figure 17: Pairwise win rate on Type A and Type B benchmarks.




Figure 18: Distribution of Type A/B scores by model.




Figure 19: Bootstrap confidence intervals on Type A/B benchmarks.




Figure 20: Raw score heatmap on Type A/B benchmarks.




Figure 21: Normalized score heatmap on Type A/B benchmarks.




Figure 22: Rank heatmap on Type A/B benchmarks.




Figure 23: Benchmark-level score spread on Type A/B benchmarks.




Figure 24: Type A and Type B grouped bar chart.




Figure 25: Type A and Type B radar plot.




Figure 26: Type A and Type B small multiples.

4 Runtime and Telemetry




Figure 27: Mean wall-clock time by model.




Figure 28: Mean generation time by model.




Figure 29: Peak CPU RAM by model.




Figure 30: Wall-clock time by benchmark.




Figure 31: Peak CPU RAM by benchmark.




Figure 32: Telemetry field coverage.