

June 10 API Results: Five-Model Comparison

1 Data Source

This report uses the frozen pre-June-11 Spanish-time snapshot in `results/api_models_pre_2026_06_11_es`. The snapshot was extracted from commit `96e05c3` (*Llava finetuning fix*), dated 2026-06-10 22:37:44 +0200.

The snapshot contains uneven benchmark coverage across the five API models. For a fair five-way comparison, this report uses the balanced overlap available for all five models: 12 Type A/Type B benchmarks. Type A corresponds to visual question answering and visual reasoning; Type B corresponds to detection and localization.

Table 1: Raw snapshot coverage before taking the balanced five-model overlap.

Model	Result files	Benchmarks	Samples	Failed samples
claude-sonnet-4-5	30	30	600	27
gemini-3.5-flash	14	14	280	44
gpt-4.1	16	16	320	95
gpt-4o	37	37	740	65
gpt-5	13	13	260	39

2 Balanced Five-Model Results

The balanced old comparison uses 12 benchmarks for every model: six Type A QA benchmarks and six Type B detection benchmarks. On this overlap, `gpt-4.1` has the highest mean score, followed closely by `gpt-4o` and `gpt-5`. `gemini-3.5-flash` is much lower in this old snapshot, consistent with the pre-fix Gemini result files.

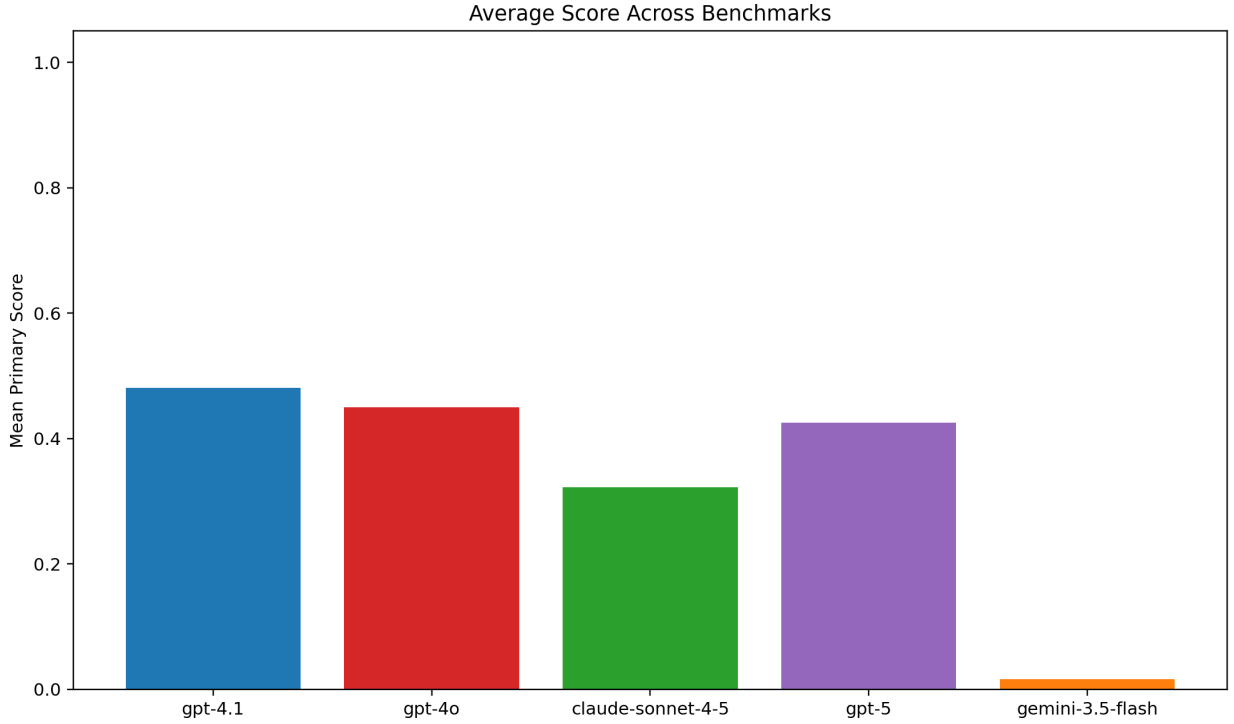
Table 2: Five-model summary across the 12 shared Type A/B benchmarks.

Model	Benchmarks	Mean score	Median score	Mean normalized
gpt-4.1	12	0.481	0.400	0.820
gpt-4o	12	0.450	0.458	0.755
gpt-5	12	0.425	0.250	0.565
claude-sonnet-4-5	12	0.322	0.293	0.606
gemini-3.5-flash	12	0.017	0.000	0.083

Table 3: Family means on the shared Type A/B overlap.

Family	Model	Benchmarks	Mean score	Mean rank
Type A / QA	gpt-4.1	6	0.867	1.17
Type A / QA	gpt-5	6	0.850	1.33
Type A / QA	gpt-4o	6	0.800	2.33
Type A / QA	claude-sonnet-4-5	6	0.508	4.00
Type A / QA	gemini-3.5-flash	6	0.033	4.83
Type B / Detection	claude-sonnet-4-5	6	0.135	1.67
Type B / Detection	gpt-4o	6	0.100	1.83
Type B / Detection	gpt-4.1	6	0.096	1.67
Type B / Detection	gpt-5	6	0.000	2.83
Type B / Detection	gemini-3.5-flash	6	0.000	2.83

3 Score and Rank Figures

**Figure 1:** Mean score for the five API models on the shared June 10 overlap.




Figure 2: Mean normalized score for the five API models.




Figure 3: First-place finishes across the 12 shared benchmarks.




Figure 4: Distribution of benchmark scores by model.




Figure 5: Bootstrap confidence intervals for mean benchmark score.




Figure 6: Raw score heatmap across models and benchmarks.




Figure 7: Normalized score heatmap across models and benchmarks.




Figure 8: Rank heatmap across models and benchmarks.




Figure 9: Benchmark-level score spread.




Figure 10: Pairwise win-rate heatmap.

4 Type A and Type B Figures




Figure 11: Grouped bar chart for Type A QA and Type B detection benchmarks.




Figure 12: Family radar plot for Type A and Type B benchmarks.

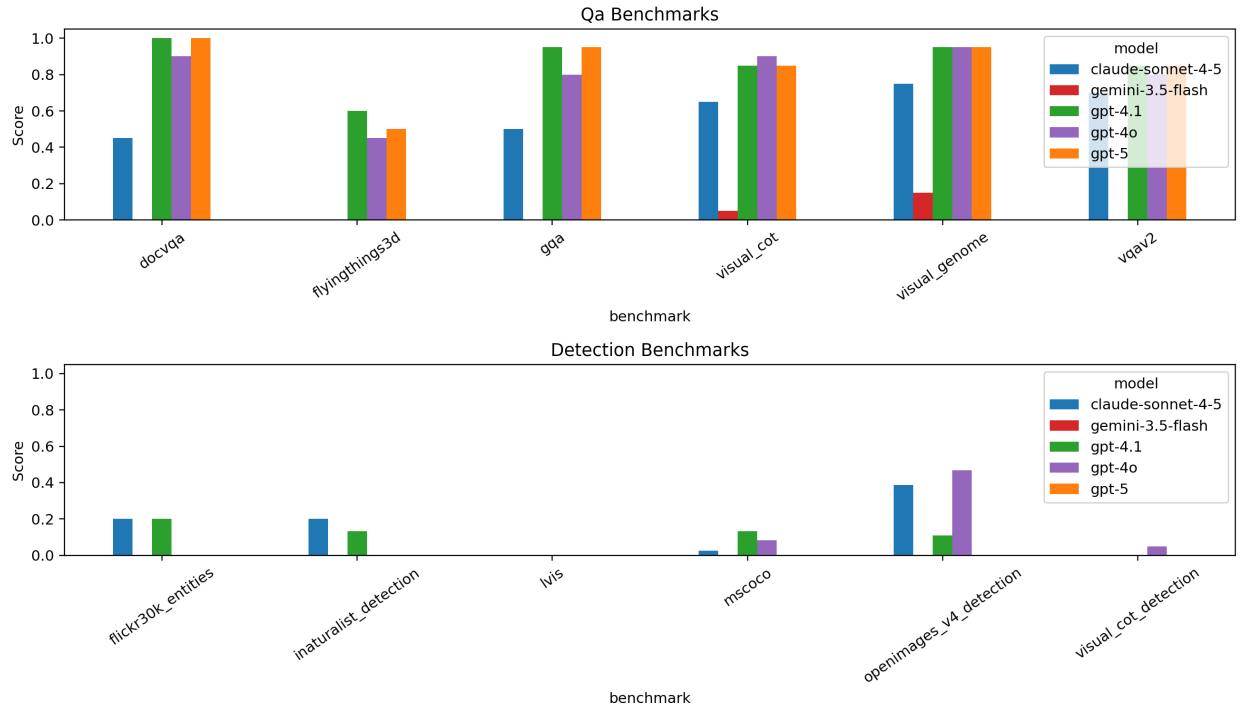


Figure 13: Small multiples for Type A and Type B benchmark families.

5 Pairwise Notes

On the balanced overlap, **gpt-4.1** beats **gpt-4o** on seven benchmarks, loses on three, and ties on two. **gpt-4o** and **gpt-5** split evenly with four wins each and four ties. Every non-Gemini model beats **gemini-3.5-flash** decisively in this old snapshot.

6 Runtime and Telemetry




Figure 14: Mean wall-clock time by model.




Figure 15: Mean generation time by model.




Figure 16: Peak CPU RAM by model.




Figure 17: Wall-clock time by benchmark.




Figure 18: Peak CPU RAM by benchmark.




Figure 19: Telemetry field coverage.