

Vision-Language Benchmarking

Short Beamer Example

Samuel

July 1, 2026

- Modern vision-language models can caption, classify, answer questions, detect objects, and compare images.
- A single leaderboard hides differences between task families.
- A benchmark harness makes heterogeneous datasets comparable under one evaluation pipeline.

Benchmark Types

Type	Task	Example dataset
C	Captioning	Conceptual Captions
A	Visual QA	FlyingThings3D
B	Detection	MS COCO
R	Rating	TAD66K
P	Preference	Pick-a-Pic

Example Input

Image A



Image B



Prompt

Which image is more aesthetically pleasing overall?

- A. Image A
- B. Image B

GPT-4o response: A

Reference: Image A

Takeaway

Main idea

The harness should record both quality and failure behavior: predictions, parsed answers, references, errors, runtime, and task-specific metrics.

This makes model comparisons auditable rather than only producing a final score.